



The Library Magic Assistant

A natural-language concierge for the catalog you already have

Prepared for **Jennifer Goodson**, Library Director, Fort Smith Public Library — for discussion.

Customers increasingly expect to ask for what they need in plain language and get an instant, accurate answer. The library already owns the hard assets — a rich catalog in **Polaris**, four branches, an events program, and licensed research databases. The Library Magic Assistant adds a friendly, natural-language **front door** to those assets: available 24/7, in any language, and welcoming to first-time researchers, caregivers, and job seekers. It extends the reference desk — it does not replace it.

What customers could do


- **Find & reserve in plain English**, with live availability across all four branches.
- **Place holds** and choose a pickup branch; join the hold queue when copies are out.
- **Request interlibrary loans** when an item isn't owned.
- **Age-appropriate suggestions** with caregiver content notes.
- **Research help** — routed to the right licensed databases, augmented with *cited* sources.
- **Discover events & story times** by age, branch, or day.

All six are clickable today in the working prototype.

Trust & governance — non-negotiables

- **Read-only replica** — the assistant physically cannot alter the live ILS.
- **No customer PII** in the AI layer; account actions use official, authenticated APIs.
- **Whitelisted functions only** — the model never writes raw database queries, which defends against prompt-injection.
- **Grounded & cited** answers — no fabricated facts.
- **Full audit log**; ambiguous or sensitive requests escalate to staff.
- The existing catalog & website stay fully usable — the assistant **augments**.

How it works



The assistant interprets the question, then calls a small set of approved functions. **Reads** hit a read-only copy of catalog and availability data. **Actions** (holds, ILL, account) go through Polaris's official authenticated web services, scoped to the customer's session. **Research** routes to licensed databases and adds cited open sources — every answer carries its citations.



Phased roadmap

Phase 1 — Read-only concierge ≈ 6-10 weeks
Hours, locations, events, "is it on the shelf?", recommendations, research routing. **No writes at all.**
Proves demand and accuracy at near-zero risk.

Phase 2 — Authenticated actions next quarter
Holds, renewals, and ILL through Polaris's official APIs; library-card sign-in; staff dashboard.

Phase 3 — Reach & depth following
Multilingual, voice/phone, kiosk mode, deeper research augmentation, and analytics to inform collection decisions.

Next steps — 30 / 60 / 90 days

- 1 **Align on scope & success metrics** with leadership — define what "good" looks like at 90 days.
- 2 **Technical discovery with Polaris / Innovative** — confirm read exports, write APIs (PAPI), and database licensing terms.
- 3 **Stand up the Phase-1 pilot** on real, read-only data behind a staff login; test with reference staff.
- 4 **Soft-launch to customers** on one high-traffic page; measure usage, accuracy, and staff time saved.
- 5 **Decide Phase 2 + hosting model** from pilot results.

Where it runs — local, cloud, or hybrid

Almost the entire system can run **on-premises**. The only piece that benefits from outside help is the heavy AI model — and even that can be local.

Component	Local?
Web app / customer UI	✓ Fully local (static site)
Read-only catalog replica	✓ Local DB fed by Polaris exports
Orchestration + guardrails	✓ Local service
Semantic search / embeddings	✓ Local (open models)
The language model	🔗 Local on a GPU box, <i>or</i> cloud API

Recommended: hybrid. Keep all data + retrieval + a local model in-building for routine questions; optionally call a cloud model for the hardest turns, sending only non-PII catalog snippets under a zero-retention agreement. Customer data stays home; answer quality stays high.

Rough cost envelope

Option	Up-front	Ongoing
Phase-1 pilot, cloud model	Integration effort	Low hundreds \$/mo at pilot volume
On-prem appliance (8-14B model + full stack)	~\$5-25k GPU server	≈ \$0 per query + staff time
Higher quality / concurrency on-prem	\$15-40k GPU	Power + maintenance

The library's existing internal stack already proves the local pattern: Postgres + vector search + a local model server. The same building blocks underpin Phase 1.

The ask: endorse a low-risk **Phase-1 pilot** on read-only data, choose a **name** for the assistant, and authorize a short technical-discovery call with the ILS vendor.